

# MA3227 Numerical Analysis II

## Lecture 08: Monte Carlo Methods

Simon Etter



Semester II, AY 2020/2021

# Monte Carlo Methods

## Problem statement

Estimate the expectation  $\mathbb{E}[X]$  of a random variable  $X$  by computing the average of a large number of samples of  $X$ .

This problem statement raises several questions.

- ▶ *Why compute expectations?*

According to basic probability theory, the expectation  $\mathbb{E}[X]$  of a random variable  $X$  which assumes values  $x$  in a discrete or continuous set  $\mathcal{X}$  with probability  $p(x)$  is given by, respectively,

$$\mathbb{E}[X] = \sum_{x \in \mathcal{X}} x p(x) \quad \text{or} \quad \mathbb{E}[X] = \int_{\mathcal{X}} x p(x) dx.$$

Computing an expectation is hence the same as evaluating a particular sum or integral, and conversely any sum or integral can be reinterpreted as an expectation.

- ▶ *Why compute expectations through sampling?*

This question is best answered by means of some examples; see the following slides.

# Monte Carlo Methods

## Example 1: Winning probabilities in $m, n, k$ -games

Consider the  $m, n, k$ -game described under

<https://en.wikipedia.org/wiki/m,n,k-game>.

Assume we want to compute the probability  $P$  of a win for player 1 assuming random moves on behalf of both players.

This probability could be computed as the sum

$$P = \sum_{\text{all possible games}} [\text{probability of game}] \times \begin{cases} 1 & \text{if player 1 wins,} \\ 0 & \text{otherwise,} \end{cases}$$

but this sum contains roughly  $(mn)!$  terms and therefore cannot be evaluated except for very modest values of  $m$  and  $n$ .

For example, if we assume a runtime of just 1 nanosecond per game, then evaluating the above sum for  $m = n = 4$  would take about 6 hours, and evaluating the sum for  $m = n = 5$  would take about 500'000 years!

# Monte Carlo Methods

## Example 1: Winning probabilities in $m, n, k$ -games (continued)

These ludicrous runtimes can be avoided if we rewrite the above sum as the expectation

$$P \approx \mathbb{E}[X] \quad \text{where} \quad X = \begin{cases} 1 & \text{if player 1 wins,} \\ 0 & \text{otherwise} \end{cases}$$

and then estimate this expectation as follows.

- ▶ Play out  $N$  random games.
- ▶ For each such game  $i$ , record in the variable  $X_i \in \{0, 1\}$  whether player 1 won.
- ▶ Estimate  $\mathbb{E}[X] \approx \frac{1}{N} \sum_{i=1}^N X_i$ .

This approach is known as *Monte Carlo sampling*, and `mnk_probabilities()` shows that it leads to reasonably accurate estimates already for moderate values of  $N$ .

# Monte Carlo Methods

## Example 2: High-dimensional integrals

Assume we want to compute a  $d$ -dimensional integral

$$I = \int_0^1 \dots \int_0^1 f(x_1, \dots, x_d) dx_1 \dots dx_d.$$

Perhaps the most obvious way to approximate this quantity is to apply a one-dimensional quadrature rule  $(x_k, w_k)_{k=1}^n$  to each of these  $d$  integrals, i.e. to compute

$$I \approx \sum_{k_d=1}^n \dots \sum_{k_1=1}^n f(x_{k_1}, \dots, x_{k_d}) w_{k_1} \dots w_{k_d}.$$

We observe:

- ▶ The above approximation requires  $N = n^d$  function evaluations.
- ▶ If the one-dimensional quadrature rule has an  $O(n^{-p})$  error, then so does the high-dimensional quadrature approximation.

*Proof.* See next slide.

# Monte Carlo Methods

*Proof that high-dimensional quadrature inherits the order of convergence of the one-dimensional quadrature rule.*

Iteratively replacing integrals with quadrature approximations, we obtain

$$\begin{aligned} & \int_0^1 \dots \int_0^1 f(x_1, \dots, x_d) dx_1 \dots dx_d = \dots \\ &= \int_0^1 \dots \int_0^1 \left( \sum_{k_1=1}^n f(x_{k_1}, x_2, \dots, x_d) w_{k_1} + O(n^{-p}) \right) dx_2 \dots dx_d \\ &= \dots \\ &= \sum_{k_d=1}^n \dots \sum_{k_1=1}^n f(x_{k_1}, \dots, x_{k_d}) w_{k_1} \dots w_{k_d} + O(n^{-p}). \end{aligned}$$

# Monte Carlo Methods

## Example 2: High-dimensional integrals (continued)

Our observations on slide 5 imply that the error and number of function evaluations  $N = n^d$  are related through

$$\text{error} = O(N^{-p/d});$$

see `integral_via_quadrature()`.

The number of function evaluations required to meet a certain error tolerance is hence given by

$$N = O(\text{error}^{-d/p}),$$

i.e.  $N$  scales exponentially in the number of dimensions  $d$  and therefore becomes prohibitively large already for moderate values of  $d$ .

This phenomenon is known as the “curse of dimensionality”.

# Monte Carlo Methods

## Example 2: High-dimensional integrals (continued)

As in the  $m, n, k$ -game example, it turns out that we can avoid excessive runtimes by replacing deterministic computations with random sampling.

In the case of high-dimensional integrals, the key to doing so is to observe that we can interpret such integrals as the expectation

$$\mathbb{E}[f(X)] = \int_0^1 \dots \int_0^1 f(x_1, \dots, x_d) dx_1 \dots dx_d$$

and hence

$$\mathbb{E}[f(X)] \approx \frac{1}{N} \sum_{k=1}^N f(X_{k,1}, \dots, X_{k,d})$$

where

- ▶  $X$  is a random variable uniformly distributed over  $[0, 1]^d$ , and
- ▶  $X_k \in [0, 1]^d$  is a sequence of samples of this random variable.


This trick is again known as *Monte Carlo sampling*.




# Monte Carlo Methods

## Example 2: High-dimensional integrals (continued)

Illustration of quadrature vs. Monte Carlo sampling:



Quadrature



Monte Carlo

Blue dots denote quadrature / sampling points.

# Monte Carlo Methods

## Example 2: High-dimensional integrals (continued)

`integral_via_monte_carlo()` shows that Monte Carlo sampling converges with rate  $O(N^{-1/2})$  regardless of the dimension.

We therefore conclude:

- ▶ Monte Carlo sampling is **slower** than midpoint quadrature if

$$2/d > 1/2 \iff d < 4.$$

- ▶ Monte Carlo sampling is **faster** than midpoint quadrature if

$$2/d < 1/2 \iff d > 4.$$

In particular, Monte Carlo sampling avoids the curse of dimensionality and is therefore a powerful tool to tackle high-dimensional problems.

# Monte Carlo Methods

## **Example 1: Winning probabilities in $m, n, k$ -games (reprise)**

It turns out that now that we know what to look out for, we can also observe the  $O(N^{-1/2})$  convergence behaviour when applying Monte Carlo sampling to compute the  $m, n, k$ -winning probabilities: every time we run `mnk_probabilities()` with a 100x larger number of samples  $N$ , we gain roughly one extra digit of accuracy.

# Monte Carlo Methods

## Introduction to Monte Carlo error estimation

The above examples indicate that the power of the Monte Carlo approach stems from the fact that

$$\mathbb{E}[X] = \frac{1}{N} \sum_{k=1}^N X_k + O(N^{-1/2})$$

for a very wide class of random variables  $X$ .

Our goal in the following will be to clarify and prove this claim. Doing so requires a solid background in probability theory; hence let us begin by recapitulating the basics.

# Monte Carlo Methods

## Probability theory in a nutshell

- ▶ A *measure* is a function  $P$  which maps subsets of some set  $\Omega$  to nonnegative real numbers such that

$$A, B \subset \Omega \quad \text{and} \quad A \cap B = \{\} \quad \implies \quad P(A \cup B) = P(A) + P(B).$$

- ▶ A *probability measure* is a measure  $P$  such that  $P(\Omega) = 1$ .
- ▶ A *probability space* is a pair  $(\Omega, P)$  such that  $P$  is a probability measure on the set  $\Omega$ .
- ▶ A *random variable* is a function  $X : \Omega \rightarrow \mathcal{X}$  defined on a probability space  $(\Omega, P)$ .
- ▶ The probability measure  $\hat{P}$  on  $\mathcal{X}$  given by

$$\hat{P}(A) = P(\{\omega \in \Omega \mid X(\omega) \in A\})$$

is called the *distribution* of the random variable  $X : \Omega \rightarrow \mathcal{X}$ .

- ▶ Probabilities associated with random variables are commonly expressed using the following convenience notations:

$$P(X \in A) = P(\{\omega \in \Omega \mid X(\omega) \in A\}),$$

$$P(X_1 \in A_1, X_2 \in A_2) = P(\{\omega \in \Omega \mid X_1(\omega) \in A_1 \wedge X_2(\omega) \in A_2\}).$$

# Monte Carlo Methods

## Example 1: Winning probabilities in $m, n, k$ -games (continued)

The above definitions were fairly abstract, so let me illustrate them by means of the  $m, n, k$ -games example from the beginning of this lecture. This example can be mapped to the above framework as follows.

- ▶ The probability space  $\Omega$  is the set of all sequences of board states

$$\omega = B_0 \rightarrow B_1 \rightarrow \dots \rightarrow B_g$$

which represent a rules-conforming and complete game.

- ▶ Since this set is discrete, we can specify the probability measure  $P$  by specifying the probability  $P(\{\omega\})$  of each sequence  $\omega \in \Omega$  and then set

$$P(A) = \sum_{\omega \in A} P(\{\omega\}) \quad \text{for all } A \subset \Omega.$$

# Monte Carlo Methods

## Example 1: Winning probabilities in $m, n, k$ -games (continued)

- ▶ According to the “random moves” assumption, the probability of a game of length  $g$  is given by

$$P(\{B_0 \rightarrow \dots \rightarrow B_g\}) = \prod_{i=0}^g (mn - i)^{-1}.$$

- ▶ The winning probability for player 1 is then given by  $P(W_1 = 1)$  where  $W_1$  denotes the random variable  $W_1 : \Omega \rightarrow \{0, 1\}$  given by

$$W_1(B_0 \rightarrow \dots \rightarrow B_g) = \begin{cases} 1 & \text{if } B_g \text{ shows a win for player 1, and} \\ 0 & \text{otherwise.} \end{cases}$$

# Monte Carlo Methods

[To be continued]